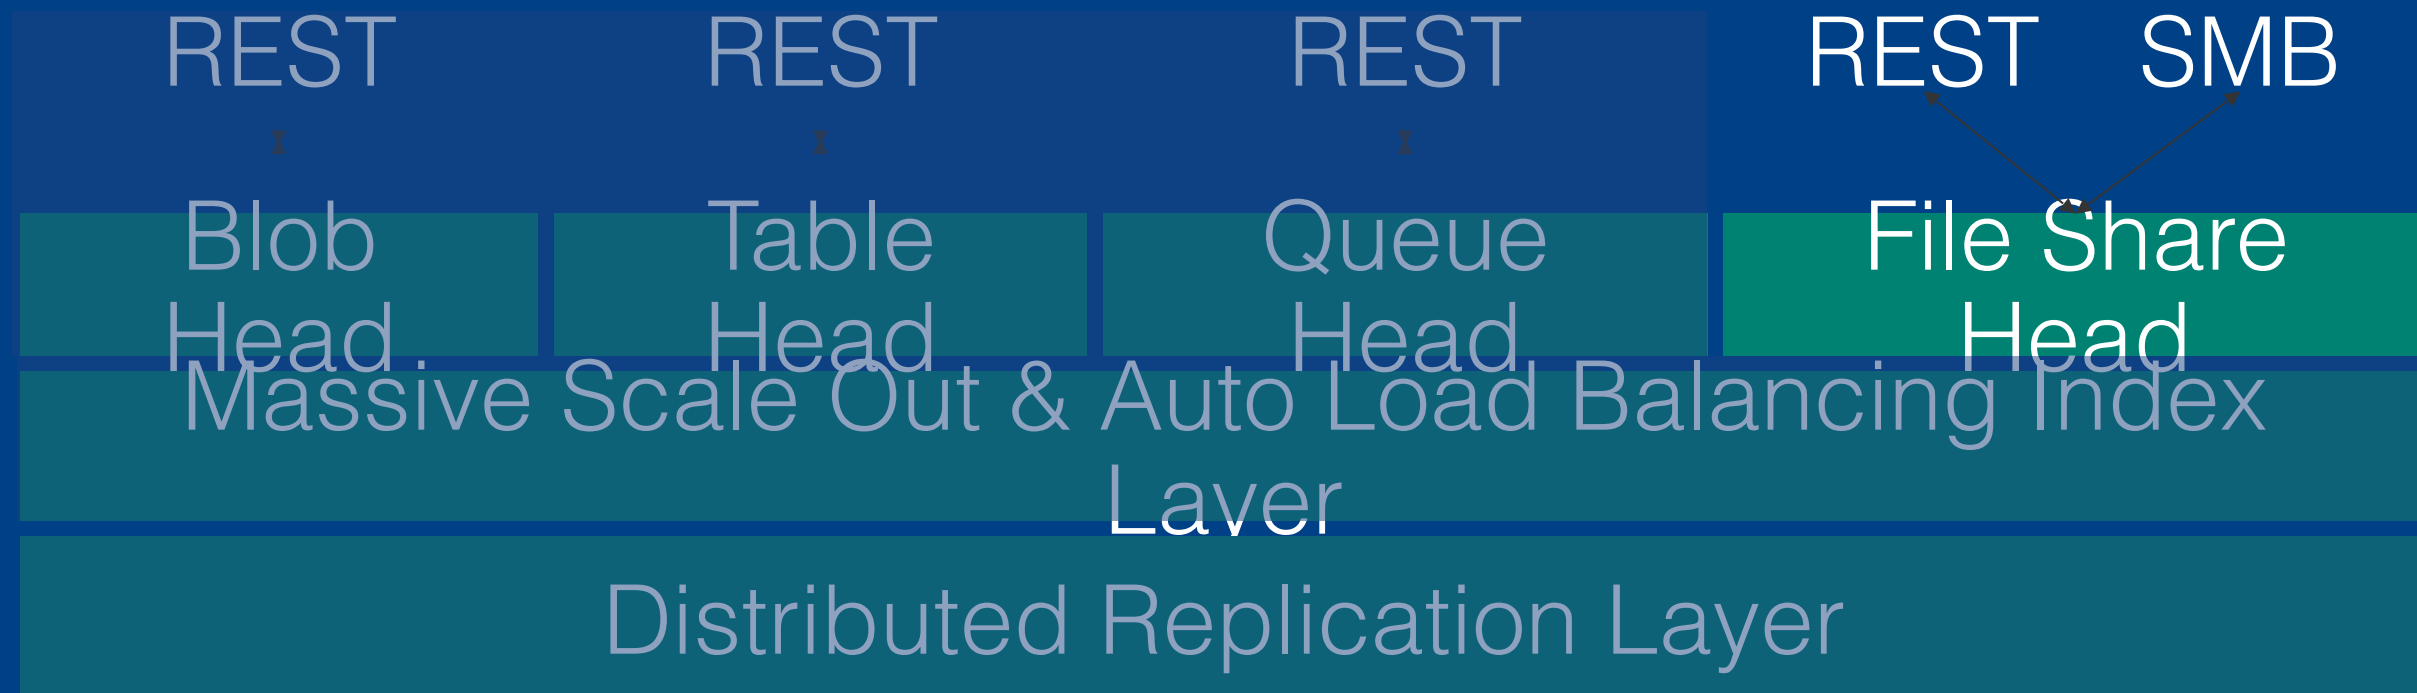


# Azure Data Storage




# Azure Storage Architecture

Microsoft Azure



# Blobs



# Microsoft Azure Storage Blob

# Two Types of Blobs Under the Hood

Block Blob

Page Blob

# Block Blob

Targeted at streaming workloads

Each blob consists of a sequence of blocks

Each block is identified by a Block ID

Size limit 200GB per blob

Optimistic Concurrency via Etags



# Page Blob

Targeted at random read/write workloads

Each blob consists of an array of pages

Each page is identified by its offset from the start of the blob

Size limit 1TB per blob

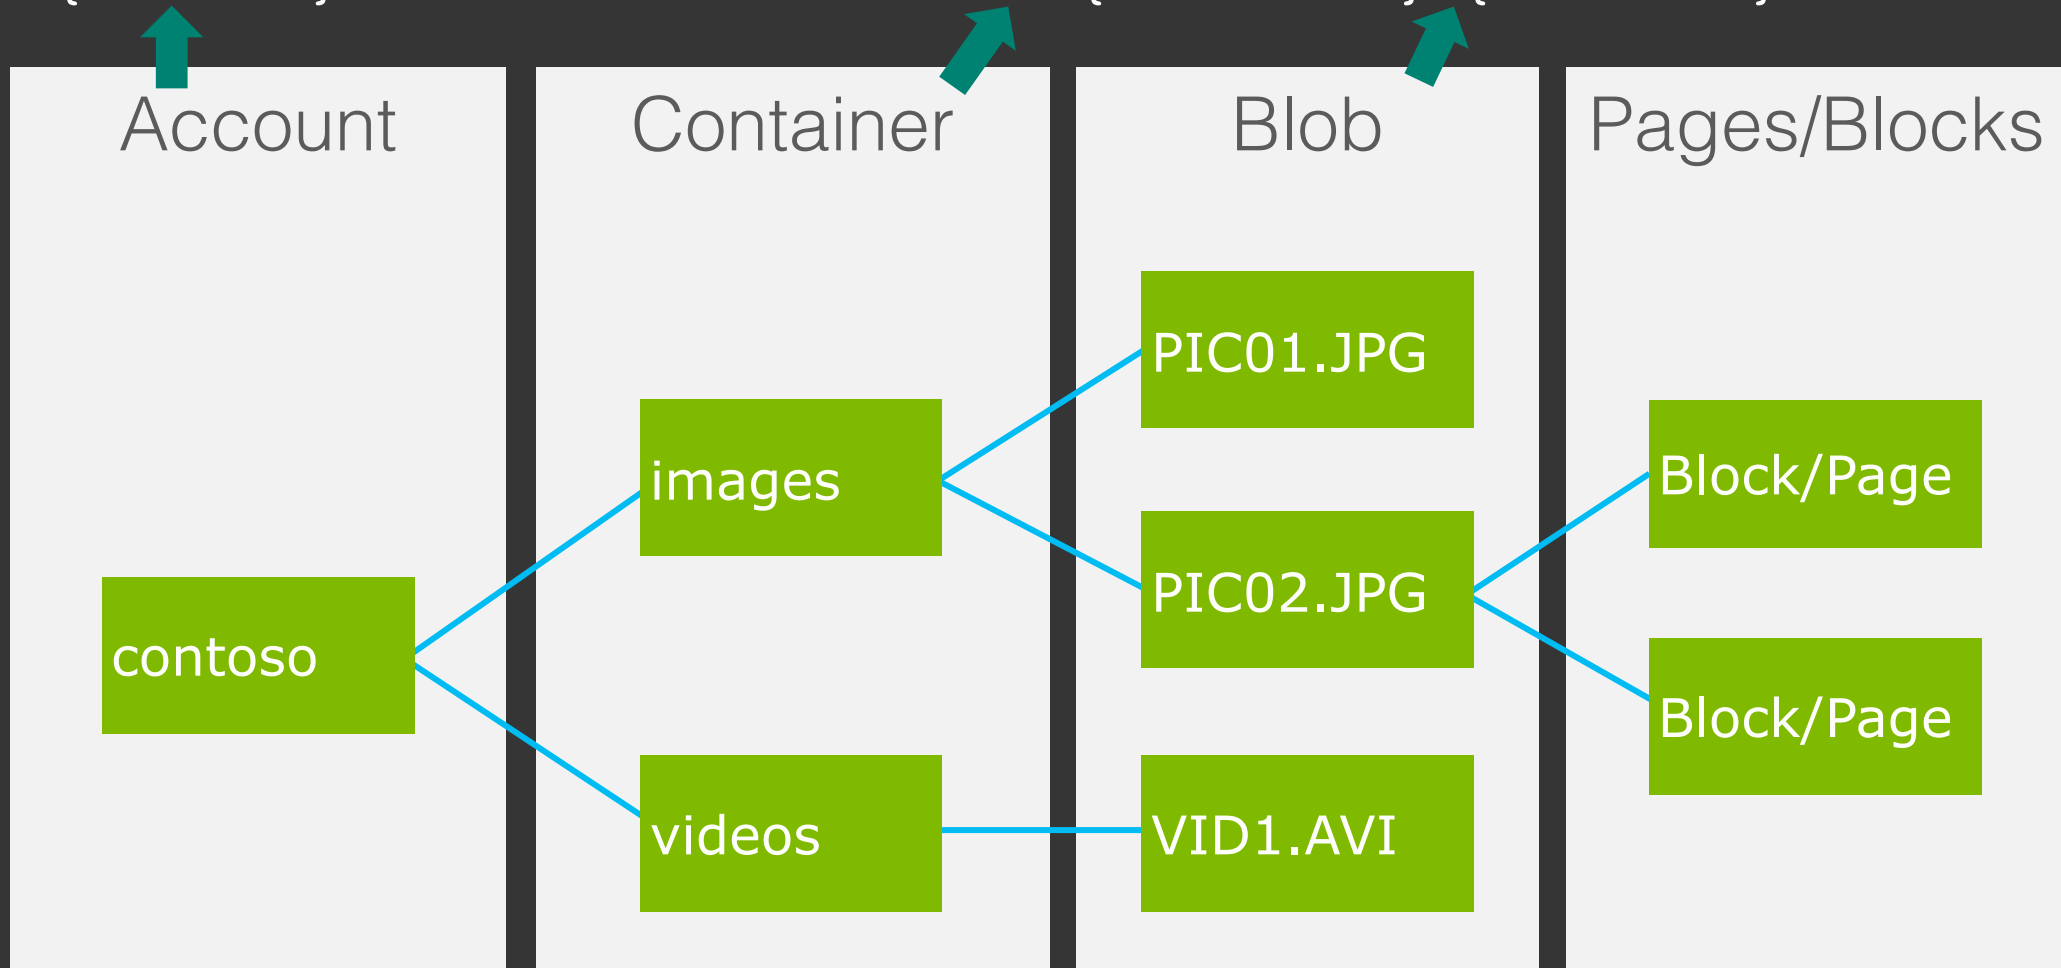
Optimistic or Pessimistic (locking) concurrency via leases



# Blob Storage Concepts

Microsoft Azure

`http://{account}.blob.core.windows.net/{container}/{blobname}`





# Blob Details – Containers

- Multiple Containers per Account
- Special \$root container



# Blob Details – Containers

- A container holds a set of blobs
- Set access policies at the container level
- Associate Metadata with Container
- List the blobs in a container
- Including Blob Metadata and MD5  
no search on metadata WHERE MetadataValue = ?



# Blob Details – Throughput

- Effectively in Partition of 1
- Target of 60MB/s per Blob



PutBlob  
GetBlob  
DeleteBlob  
CopyBlob  
SnapshotBlob  
LeaseBlob



# Blob Details

## Associate metadata with blob


- Standard HTTP metadata/headers (Cache-Control, Content-Encoding, Content-Type, etc)
- Metadata is <name, value> pairs, up to 8KB per blob
- Either as part of PutBlob or independently

# Blob Details – Blob always accessed by name

Can include '/' or other delimiter in name

e.g. /<container>/myblobs/smurf.png

  
blob name



GET Blob operation takes parameters

Prefix

Delimiter

Include = (snapshots,  
metadata etc...)

Blob Details



# Blob sample listing

```
http://  
adventureworks.blob.core.windows.net/  
    Products/Bikes/  
SuperDuperCycle.jpg  
    Products/Bikes/FastBike.jpg  
    Products/Canoes/Hybrid.jpg  
    Products/Canoes/Flatwater.jpg  
    Products/Canoes/Whitewater.jpg  
    Products/Tents/PalaceTent.jpg  
    Products/Tents/ShedTent.jpg
```

GET http://.../products?comp=list&prefix=Tents

```
<Blobs>  
  <Blob><Name>Tents/PalaceTent.jpg</Name>[...]</  
Blob>  
  <Blob><Name>Tents/ShedTent.jpg</Name>[...]</
```



# Blob sample listing full response

```
<Blobs>
  <Blob>
    <Name>Tents/PalaceTent.jpg</Name>
    <Url>https://readinesscloudcamp.blob.core.windows.net/products/Tents/PalaceTent.jpg</Url>
    <LastModified>Wed, 17 Dec 2014 09:00:26 GMT</LastModified>
    <Etag>0x8D1E7EF08F31520</Etag>
    <Size>150027</Size>
    <ContentType>image/jpeg</ContentType>
    <ContentEncoding />
    <ContentLanguage />
  </Blob>
  <Blob>
    <Name>Tents/ShedTent.jpg</Name>
    <Url>https://readinesscloudcamp.blob.core.windows.net/products/Tents/ShedTent.jpg</Url>
    <LastModified>Wed, 17 Dec 2014 09:00:26 GMT</LastModified>
    <Etag>0x8D1E7EF08EA6257</Etag>
    <Size>150027</Size>
    <ContentType>image/jpeg</ContentType>
    <ContentEncoding />
    <ContentLanguage />
  </Blob>
</Blobs>
```

# Blob sample listing with maxresults

```
http://  
adventureworks.blob.core.windows.net/  
    Products/Bikes/  
SuperDuperCycle.jpg  
    Products/Bikes/FastBike.jpg  
    Products/Canoes/Hybrid.jpg  
    Products/Canoes/Flatwater.jpg  
    Products/Canoes/Whitewater.jpg  
    Products/Tents/PalaceTent.jpg  
    Products/Tents/ShedTent.jpg  
http://.../products?
```

comp=list&prefix=Canoes&maxresults=2

```
<Blob>Canoes/Hybrid.jpg</Blob>  
<Blob>Canoes/Flatwater.jpg</Blob>  
<NextMarker>1!28!Q2Fub2VzL1doaXRld2F0ZXluanBn</  
NextMarker>
```

# Blob sample listing with maxresults

```
http://  
adventureworks.blob.core.windows.net/  
    Products/Bikes/  
SuperDuperCycle.jpg  
    Products/Bikes/FastBike.jpg  
    Products/Canoes/Hybrid.jpg  
    Products/Canoes/Flatwater.jpg  
    Products/Canoes/Whitewater.jpg  
    Products/Tents/PalaceTent.jpg  
    Products/Tents/ShedTent.jpg  
http://.../products?
```

```
    comp=list&prefix=Canoes&maxresults=2  
&marker=1!28!Q2Fub2VzL1doaXRld2F0ZXIuanBn  
    <Blob>Canoes/Whitewater.jpg</  
    Blob>  
    </NextMarker>
```

# Uploading a Block Blob

## Uploading



## THE BLOB

```
blobName = "TheBlob.wmv";  
PutBlock(blobName, blockId1,  
block1Bits);  
PutBlock(blobName, blockId2,  
block2Bits);  
.....  
PutBlock(blobName, blockIdN,  
blockNBits);  
PutBlockList(blobName,  
blockId1,...,blockIdN);
```

TheBlob.wm  
v

# Blob block uploading benefits

Efficient continuation and retry

Parallel and out of order upload of blocks



# Page Blob – Random Read/Write

Create blob and specify Blob Size = 10 Gbytes

Fixed Page Size = 512 bytes

Random Access Operations:

`PutPage[512, 2048)`

`PutPage[0, 1024)`

`ClearPage[512, 1536)`

`PutPage[2048, 2560)`


`GetPageRange[0, 4096)` returns valid data ranges:

`[0, 512)` , `[1536, 2560)`


`GetBlob[1000, 2048)` returns:

All 0 for first 536 bytes

Next 512 bytes data stored in `[1536, 2048)`  Microsoft



# Page Blob – Random Read/Write



Sparse storage:  
Only charged for pages with data  
stored in them

## Shared Access Signatures

Fine grain access rights to blobs and containers  
Sign URL with storage key – permit elevated rights





# Shared Access Signatures – Two broad approaches

Ad-hoc:  
Stored Access Policy

Policy based:  
Shared Access Signature

# Shared Access Signatures – Revocation

Use short time periods and re-issue

Use container level policy that can be deleted



# Shared Access Signatures – Ad Hoc Signatures

## Create Short Dated Shared Access Signature

Signed resource Blob or Container

AccessPolicy Start, Expiry and Permissions

Signature HMAC-SHA256 of above fields



# Shared Access Signatures – Ad Hoc Signatures

## Use case

Single use URLs

E.g. Provide URL to mobile client to upload to container



```
http://...blob.../pics/image.jpg?  
sr=c&st=2009-02-09T08:20Z&se=2009-02-10T08:30Z&sp=w  
&sig= dD80ihBh5jfNpym05Hg1IdiJIEvHcJpCMiCMnN%2fRnbI%3d
```

# Store Access Policy – Policy Based Signatures

## Create Container Level Policy

Specify StartTime, ExpiryTime, Permissions



# Store Access Policy – Policy Based Signatures

## Create Shared Access Signature URL

Signed resource Blob or Container

Signed identifier Optional pointer to container policy

Signature HMAC-SHA256 of above fields



# Store Access Policy – Policy Based Signatures

## Use case

Providing revocable permissions to certain users/groups

To revoke: Delete or update container policy





```
http://...blob.../pics/image.jpg?  
sr=c&si=MyUploadPolicyForUserID12345  
&sig=dD80ihBh5jfNpymO5Hg1IdiJIEvHcJpCMiCMnN%2fRnbI%3d
```

# Files

“I wish I could go to storage and provision a cloud drive, giving it a namespace, and that drive would then be UNC-addressable by the OSes.”

“I need two VM's running with a shared drive. One will write to the drive, the other will read [it].”

“Hi, I have two VM's in Microsoft Azure. All I want to do is set up a file share between them. Is this possible?”

“Is it possible to share a secondary disk between different VM instances?”

Setup an IaaS VM to host a File Share backed by an IaaS Disk

Write code to find the IaaS File Share from the rest of the VMs



Write some code to provide high availability

Handle (Sharing IaaS Disk) the fa

Backup IaaS VMs  
(Mount/Share after failover)


You can only access the File Share from other VMs

Sharing Files –  
The old way



# Azure Files

- Shared Network File Storage for Azure
- Availability, durability, scalability are managed automatically
- Supports two interfaces: SMB and REST



# Azure Files – Usage

- Share data across VMs and applications
- Share settings throughout services
- Dev/Test/Debug

# Queues

# Why use a Queue?

Queue length reflects how well the backend processing nodes are doing.

Decouples the application.

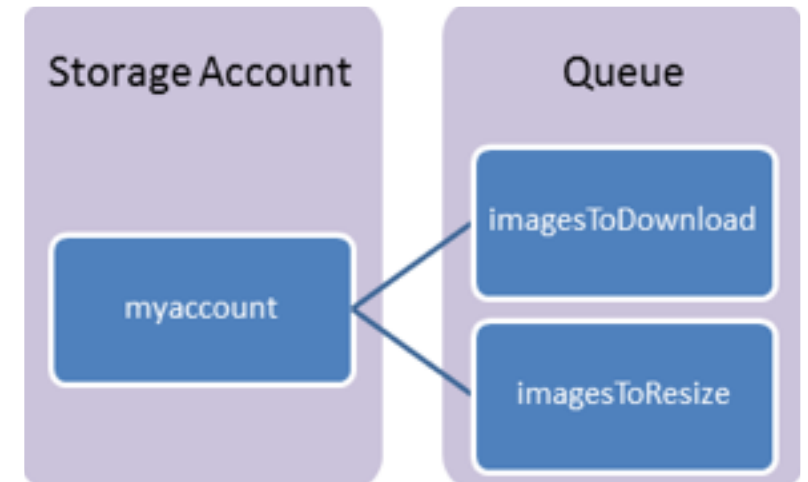
Flexibility of efficient resource usage within an application.

Absorb traffic bursts and reduce the impact of individual component failures.



# Queue Components

- Storage Account: All access to Azure Storage is done through a storage account.
- Queue: A queue contains a set of messages.
- Message: A message, in any format, of up to 64KB.



# Queue URL format

Queues are addressable using the following URL format:


`http://{storage-account}.queue.core.windows.net/{queue}`

# Queue URL format

Example:

<http://myaccount.queue.core.windows.net/imagesToDownload>

# Queue-based Load Levelling Pattern



Messages are ordered but not guaranteed FIFO.

Message will be processed at least once.

Message may be processed more than once.

.DequeueCount increases every time.


-> Processing must be idempotent.

## Queue Considerations

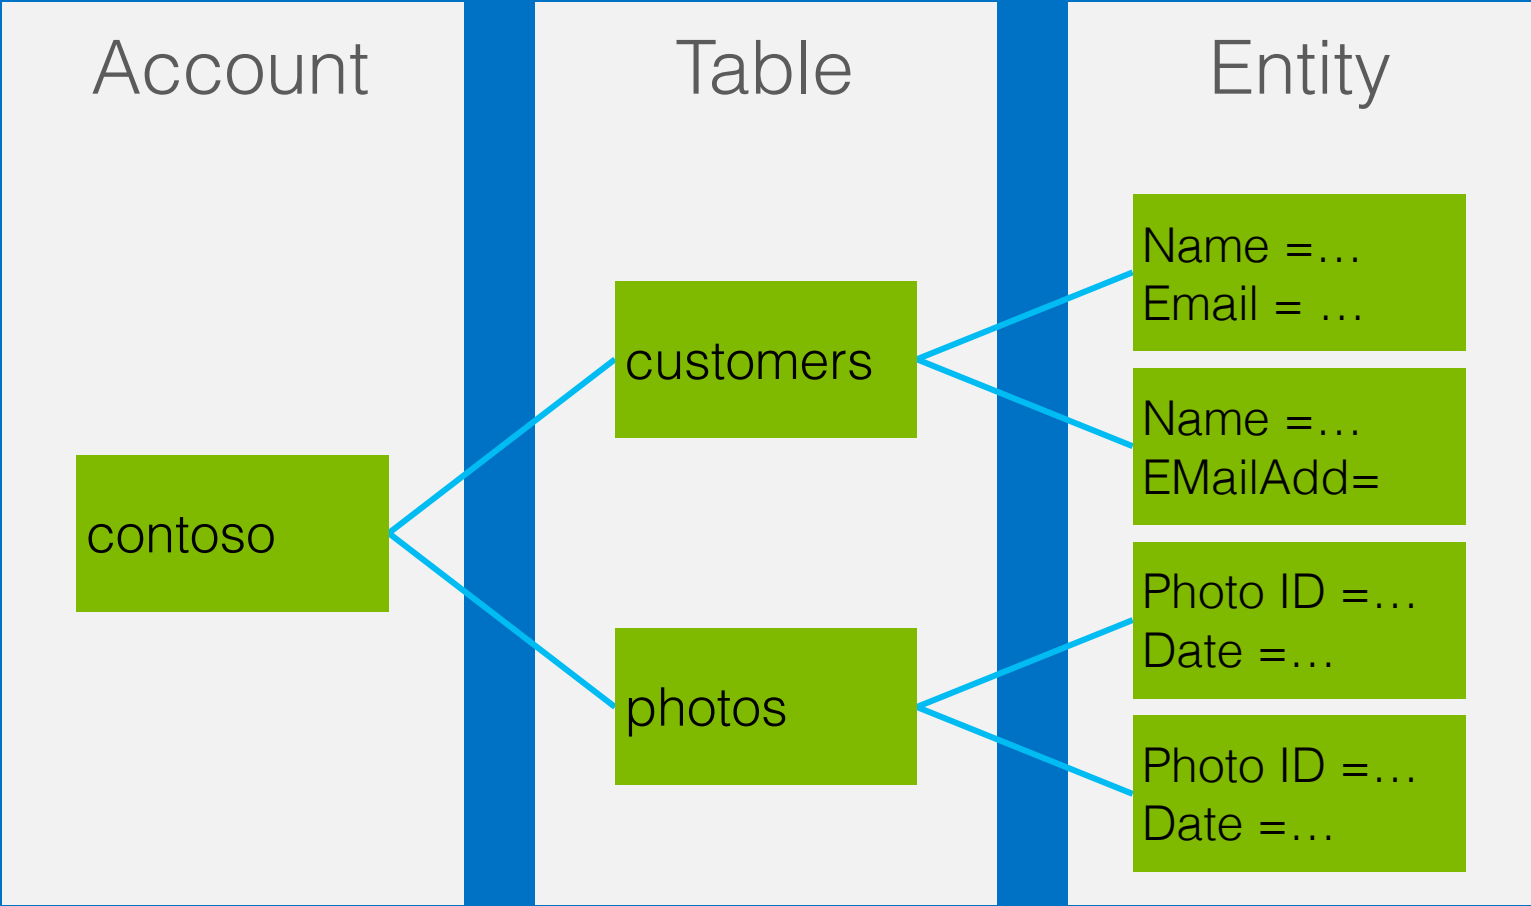


# Queue Considerations

Messages are stored up to 7 days



# Tables





# Table Storage Details

Not an RDBMS Table!  
The mental picture is 'Entities'



# Table Storage Details

Entity can have up to 255 properties

Up to 1MB per entity



# Table Storage Details

## Entity Properties

PartitionKey & RowKey are mandatory properties

Composite key which uniquely identifies an entity

They are the only indexed properties

Defines the sort order

# Table Storage Details

## Purpose of the PartitionKey

**Entity Locality**

**Entity Group Transactions**

**Table Scalability**

# Table Storage Details

## Purpose of the PartitionKey

### **Entity Locality**

Entities in the same partition will be stored together

Efficient querying and cache locality

Endeavour to include partition key in all queries

# Table Storage Details

## Purpose of the PartitionKey

### **Entity Group Transactions**

Atomic multiple CRUD in same partition in a single transaction

# Table Storage Details

## Purpose of the PartitionKey

### **Table Scalability**

Target throughput – 500 tps/partition, several thousand tps/account

Microsoft Azure monitors the usage patterns of partitions

# Table Storage Details

## Purpose of the PartitionKey

### **Table Scalability**

Automatically load balance partitions

Each partition can be served by a different storage node

Scale to meet the traffic needs of your table



Timestamp property

Optimistic Concurrency

Exposed as an HTTP Etag

Table Storage Details  
Entity Properties



No fixed schema for other properties

Each property is stored as: <name, typed value>

Properties can be the standard .NET types:

string, binary, bool, DateTime, GUID, int, int64, double

Table Storage Details  
Entity Properties



# Table Storage Details

Supports full manipulation (CRUD)


Including Upsert and Entity Group Transactions

Tables can have metadata



# StorSimple

# StorSimple + Microsoft Azure = Hybrid Cloud Storage



# StorSimple

## Designed to:

Reduce storage costs

Simplify storage management

Improve disaster recovery capability and efficiency

Provide data mobility.

